

Machine Learning Analysis Report

Iris Species Classification using Random Forest:

A Supervised Machine Learning Workflow

Udacity MSc Artificial Intelligence · Machine Learning Foundations · March 2026
Qeis Kamran · OTH Amberg-Weiden · CEREBRA Laboratories AG

1. Overview

This report presents a complete supervised machine learning workflow applied to the UCI Iris dataset (Fisher, 1936), available at <https://archive.ics.uci.edu/ml/datasets/iris>. The task is multi-class classification: predicting the species of iris flowers (Setosa, Versicolor, or Virginica) from four morphological measurements. A Random Forest classifier was selected, trained, and evaluated using accuracy, F1-score, and 5-fold cross-validation.

The Iris dataset is a canonical benchmark in pattern recognition and machine learning (Breiman, 2001). While simple by modern standards, it is well-suited for demonstrating the full supervised learning workflow — from preprocessing through evaluation — and provides clear, interpretable results that support meaningful analysis of model behaviour.

2. Dataset Description

The UCI Iris dataset contains 150 observations of iris flowers collected by Ronald Fisher in 1936. Each observation represents one flower sample described by four continuous morphological features and a categorical class label. The dataset is perfectly balanced with 50 samples per species.

Variable	Type	Description	Range
sepal length (cm)	Numeric input	Length of the sepal	4.3 – 7.9
sepal width (cm)	Numeric input	Width of the sepal	2.0 – 4.4
petal length (cm)	Numeric input	Length of the petal (most discriminative)	1.0 – 6.9
petal width (cm)	Numeric input	Width of the petal (most discriminative)	0.1 – 2.5
species	Categorical target	Iris-setosa / Iris-versicolor / Iris-virginica	3 classes

There are no missing values. Setosa is linearly separable from the other two species in petal space, while Versicolor and Virginica show partial overlap, making this a moderately challenging three-class problem (Figure 3A).

3. Modeling Approach

3.1 Data Preprocessing

The dataset was split into training (80%, n=120) and test (20%, n=30) sets using stratified sampling to preserve class proportions. Features were standardised to zero mean and unit variance using StandardScaler before training. Standardisation is recommended for ensemble tree methods to ensure that features measured on different scales do not dominate importance calculations (Pedregosa et al., 2011).

3.2 Model Selection

A Random Forest classifier (Breiman, 2001) was selected for the following reasons:

- It handles multi-class classification natively without one-vs-rest decomposition
- It is robust to the small sample size (n=150) through bootstrap aggregation
- It provides interpretable feature importance scores, directly relevant to understanding which measurements drive species discrimination
- It has low sensitivity to hyperparameter choices, making it appropriate for a foundational ML project

3.3 Evaluation Metrics

Three evaluation approaches were used. Test set accuracy measures overall correct classification rate on held-out data. The classification report provides per-class precision, recall, and F1-score, which are essential for understanding performance on each species individually. Five-fold cross-validation provides a more reliable estimate of generalisation performance by averaging over five independent train/test splits (Pedregosa et al., 2011).

4. Results

The Random Forest classifier achieved the following performance on the held-out test set:

Metric	Setosa	Versicolor	Virginica	Overall
Precision	1.00	0.82	0.89	0.90
Recall	1.00	0.90	0.80	0.90
F1-Score	1.00	0.86	0.84	0.90
Test Accuracy	—	—	—	0.9000 (90.0%)
CV Accuracy (5-fold)	—	—	—	0.9667 ± 0.0211

All 10 setosa test samples were classified correctly (Figure 2A). Three errors occurred: two versicolor samples were predicted as virginica and one virginica was predicted as versicolor. This pattern of confusion between the two non-setosa species is consistent with their overlapping petal distributions (Figure 3A).

Feature importance analysis (Figure 2B) confirmed that petal length (importance ≈ 0.44) and petal width (importance ≈ 0.42) together account for over 85% of the model's predictive power. Sepal length and sepal width contribute less than 15% combined.

The cross-validation results (Figure 3B) show consistently high accuracy across all five folds (range: 0.933 – 1.000), confirming that the model generalises well and is not overfitting to the specific training split used for the primary evaluation.

5. Interpretation for a Non-Technical Audience

In plain language: the model learned to distinguish iris flower species from measurements of their petals and sepals with 90% accuracy on new flowers it had never seen. If you showed this model 100 iris flowers, it would correctly identify the species of 90 of them.

The most important finding is that petal size — both length and width — is far more useful for species identification than sepal size. Setosa flowers are easy to identify because their petals are distinctly smaller than the other two species. Versicolor and Virginica are more similar, which is why the model occasionally confuses them.

The cross-validation result of 96.7% across five independent tests is more reliable than the single test result of 90%, suggesting the model's true performance is closer to the higher figure. The gap between 90% and 96.7% reflects the natural variability when working with only 30 test samples.

6. Limitations and Potential Bias

6.1 Dataset Limitations

The dataset contains only 150 samples from a single geographic location collected in 1936. It is one of the most studied datasets in machine learning, which means that published benchmark results provide unreliable guidance — many reported accuracies reflect overfitting to this specific dataset across the literature rather than genuine generalisation (Breiman, 2001).

The small test set (n=30) means that each misclassification changes accuracy by 3.3 percentage points. The difference between the test accuracy (90%) and cross-validated accuracy (96.7%) reflects this variance rather than genuine model instability.

6.2 Potential Sources of Bias

The perfectly balanced class distribution (50 per species) is an idealised training condition that rarely holds in real-world species classification problems. If deployed in a setting where one species is far more common than others, the model's precision and recall would differ substantially from the results reported here.

Feature importance scores from Random Forest reflect the model's internal decision structure, not the causal relationships between morphology and species identity. The finding that petal length is the most 'important' feature does not mean that petal length causally determines species — it means that petal length is the measurement that most reliably splits the training data into correct groups.

7. Workflow Completeness and Professional Relevance

This project demonstrates a complete, professional machine learning workflow that forms the direct foundation for more advanced AI work. The workflow covers every stage of applied supervised learning: problem framing, data loading and inspection, preprocessing (stratified splitting, feature scaling), model selection with justification, training, multi-metric evaluation (accuracy, F1, confusion matrix, cross-validation), and written interpretation for both technical and non-technical audiences.

The specific skills developed in this project connect directly to more advanced capstone modules:

This Project	Foundation For
StandardScaler preprocessing	Neural network input normalisation (Deep Learning module)
Random Forest feature importance	Attention mechanisms and feature attribution in Transformers
Cross-validation methodology	Hyperparameter tuning and model selection in deep learning
Confusion matrix analysis	Per-class evaluation in CNN/RNN image and sequence models
scikit-learn pipeline structure	PyTorch DataLoader and training loop architecture
Bias and class imbalance discussion	Responsible AI and fairness in Generative AI module

In my own applied AI work at CEREBRA Laboratories AG, the same workflow structure — preprocessing, model selection, evaluation, interpretability, and bias analysis — forms the core of the CEREBRA-CogAssess Bayesian Adaptive Cognitive Assessment Agent, which applies machine learning to early dementia detection. The supervised classification skills developed here are directly applicable to clinical risk stratification models where precision, recall, and class imbalance handling are safety-critical.

The notebook (modeling.ipynb) is structured to be immediately reusable: modular functions with docstrings, a

reproducible random seed, a complete requirements.txt, and version-controlled code on GitHub. Any collaborator or reviewer can clone the repository, install dependencies, and reproduce every result in the notebook within minutes.

References

Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32. <https://doi.org/10.1023/A:1010933404324>

Fisher, R. A. (1936). The use of multiple measurements in taxonomic problems. *Annals of Eugenics*, 7(2), 179–188. <https://doi.org/10.1111/j.1469-1809.1936.tb02137.x>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. <https://jmlr.org/papers/v12/pedregosa11a.html>

UCI Machine Learning Repository. (1988). Iris Data Set. University of California, Irvine. <https://archive.ics.uci.edu/ml/datasets/iris>